

✓ Convolutional Neural Networks - Build Model

In this notebook, we build and train a **CNN** to classify images from the CIFAR-10 database.

- The code provided here are **almost** working. You are required to build up a CNN model and train it.
- Make sure you covered implementations of the **TODOs** in this notebook

The images in this database are small color images that fall into one of ten classes; some example images are pictured below.



✓ Optional: Use [CUDA](#) if Available

Since these are color (32x32x3) images, it may prove useful to speed up your training time by using a GPU. CUDA is a parallel computing platform and CUDA Tensors are the same as typical Tensors, but they utilize GPU's for efficient parallel computation.

```
import torch
import numpy as np

# check if CUDA is available
train_on_gpu = torch.cuda.is_available()

if not train_on_gpu:
    print('CUDA is not available. Training on CPU ...')
else:
    print('CUDA is available! Training on GPU ...')
```

→ CUDA is available! Training on GPU ...

✓ Load the [Data](#)

Downloading may take a minute. We load in the training and test data, split the training data into a training and validation set, then create DataLoaders for each of these sets of data.

```
from torchvision import datasets
import torchvision.transforms as transforms
from torch.utils.data.sampler import SubsetRandomSampler

# number of subprocesses to use for data loading
num_workers = 0
# how many samples per batch to load
batch_size = 20
# percentage of training set to use as validation
valid_size = 0.2

# convert data to a normalized torch.FloatTensor
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# choose the training and test datasets
train_data = datasets.CIFAR10('data', train=True,
                              download=True, transform=transform)
test_data = datasets.CIFAR10('data', train=False,
                             download=True, transform=transform)

# obtain training indices that will be used for validation
num_train = len(train_data)
indices = list(range(num_train))
np.random.shuffle(indices)
split = int(np.floor(valid_size * num_train))
train_idx, valid_idx = indices[split:], indices[:split]

# define samplers for obtaining training and validation batches
train_sampler = SubsetRandomSampler(train_idx)
valid_sampler = SubsetRandomSampler(valid_idx)
```

```
# prepare data loaders (combine dataset and sampler)
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                           sampler=train_sampler, num_workers=num_workers)
valid_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                           sampler=valid_sampler, num_workers=num_workers)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                           num_workers=num_workers)

# specify the image classes
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer',
          'dog', 'frog', 'horse', 'ship', 'truck']
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to data/cifar-10-python.tar.gz
 100%|██████████| 170M/170M [00:13<00:00, 13.0MB/s]
 Extracting data/cifar-10-python.tar.gz to data
 Files already downloaded and verified

Visualize a Batch of Training Data

```
import matplotlib.pyplot as plt
%matplotlib inline

# helper function to un-normalize and display an image
def imshow(img):
    img = img / 2 + 0.5 # unnormalize
    plt.imshow(np.transpose(img, (1, 2, 0))) # convert from Tensor image

# obtain one batch of training images
dataiter = iter(train_loader)
#images, labels = dataiter.next() #python, torchvision version match issue
images, labels = next(dataiter)
images = images.numpy() # convert images to numpy for display

# plot the images in the batch, along with the corresponding labels
fig = plt.figure(figsize=(25, 4))
# display 20 images
for idx in np.arange(20):
    ax = fig.add_subplot(2, int(20/2), idx+1, xticks=[], yticks=[])
    imshow(images[idx])
    ax.set_title(classes[labels[idx]])
```






View an Image in More Detail

Here, we look at the normalized red, green, and blue (RGB) color channels as three separate, grayscale intensity images.

```
rgb_img = np.squeeze(images[3])
channels = ['red channel', 'green channel', 'blue channel']

fig = plt.figure(figsize = (36, 36))
for idx in np.arange(rgb_img.shape[0]):
    ax = fig.add_subplot(1, 3, idx + 1)
    img = rgb_img[idx]
    ax.imshow(img, cmap='gray')
    ax.set_title(channels[idx])
    width, height = img.shape
    thresh = img.max()/2.5
    for x in range(width):
        for y in range(height):
```

```
val = round(img[x][y],2) if img[x][y] !=0 else 0
ax.annotate(str(val), xy=(y,x),
            horizontalalignment='center',
            verticalalignment='center', size=8,
            color='white' if img[x][y]<thresh else 'black')
```



✓ TODO: Update the Network [Architecture](#)

Build up your own Convolutional Neural Network using Pytorch API:

- nn.Conv2d(): for convolution
- nn.MaxPool2d(): for maxpooling (spatial resolution reduction)
- nn.Linear(): for last 1 or 2 layers of fully connected layer before the output layer.
- nn.Dropout(): optional, [dropout](#) can be used to avoid overfitting.
- F.relu(): Use ReLU as the activation function for all the hidden layers

The following is a skeleton example that's not completely working.

```
import torch.nn as nn
import torch.nn.functional as F
```

Define the CNN architecture

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
```

Convolutional Layers (Extract features from the input image)

```
self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=1, padding=1) # (32x32x3) -> (32x32x32)
self.conv2 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1, padding=1) # (32x32x32) -> (32x32x64)
self.conv3 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=1, padding=1) # (32x32x64) -> (32x32x128)
```

Max Pooling Layer (Reduce spatial dimensions)

```
self.pool = nn.MaxPool2d(2, 2) # Halves the spatial size
```

Fully Connected Layers

```
self.fc1 = nn.Linear(128 * 4 * 4, 512) # Flattened size after pooling
self.fc2 = nn.Linear(512, 10) # Output layer (for 10 classes)
```

Dropout Layer (To prevent overfitting)

```
self.dropout = nn.Dropout(0.25)
```

def forward(self, x):

```
# Convolutional + ReLU + Pooling
x = self.pool(F.relu(self.conv1(x)))
x = self.pool(F.relu(self.conv2(x)))
x = self.pool(F.relu(self.conv3(x)))
```

Flatten before fully connected layers

```
x = x.view(-1, 128 * 4 * 4)
```

```

        # Fully connected layers with ReLU activation
        x = F.relu(self.fc1(x))
        x = self.dropout(x) # Dropout for regularization
        x = self.fc2(x) # Output layer

    return x

# Create a complete CNN
model = Net()
print(model)

# Move tensors to GPU if CUDA is available
if train_on_gpu:
    model.cuda()

Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=2048, out_features=512, bias=True)
  (fc2): Linear(in_features=512, out_features=10, bias=True)
  (dropout): Dropout(p=0.25, inplace=False)
)
```

▼ Specify [Loss Function](#) and [Optimizer](#)

Decide on a loss and optimization function that is best suited for this classification task. The linked code examples from above, may be a good starting point; [this PyTorch classification example](#) Pay close attention to the value for **learning rate** as this value determines how your model converges to a small error.

The following is working code, but you can make your own adjustments.

TODO: try to compare with ADAM optimizer

```

import torch.optim as optim
import torch.nn as nn
import torch.nn.functional as F

# specify loss function (categorical cross-entropy)
criterion = nn.CrossEntropyLoss()

# specify optimizer
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9) # Original choice
#optimizer = optim.Adam(model.parameters(), lr=0.001)
import torch.optim as optim

# Specify loss function (Categorical Cross-Entropy)
criterion = nn.CrossEntropyLoss()

# Optimizer Variations
#optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
#optimizer_adam = optim.Adam(model.parameters(), lr=0.001)
```

▼ Train the Network

Remember to look at how the training and validation loss decreases over time; if the validation loss ever increases it indicates possible overfitting.

The following is working code, but you are encouraged to make your own adjustments and enhance the implementation.

```

# number of epochs to train the model, you decide the number
n_epochs = 10

valid_loss_min = np.inf # track change in validation loss

# Lists to store losses for plotting
train_losses = []
valid_losses = []
```

```

for epoch in range(1, n_epochs+1):

    # keep track of training and validation loss
    train_loss = 0.0
    valid_loss = 0.0

    #####
    # train the model #
    #####
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        # move tensors to GPU if CUDA is available
        if train_on_gpu:
            data, target = data.cuda(), target.cuda()
        # clear the gradients of all optimized variables
        optimizer.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the batch loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer.step()
        # update training loss
        train_loss += loss.item()*data.size(0)

    #####
    # validate the model #
    #####
    model.eval()
    for batch_idx, (data, target) in enumerate(valid_loader):
        # move tensors to GPU if CUDA is available
        if train_on_gpu:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the batch loss
        loss = criterion(output, target)
        # update average validation loss
        valid_loss += loss.item()*data.size(0)

    # calculate average losses
    train_loss = train_loss/len(train_loader.sampler)
    valid_loss = valid_loss/len(valid_loader.sampler)

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch, train_loss, valid_loss))

    # save model if validation loss has decreased
    if valid_loss <= valid_loss_min:
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
            valid_loss_min,
            valid_loss))
        torch.save(model.state_dict(), 'model_trained_10ep_sgd.pt')
        valid_loss_min = valid_loss
    train_losses.append(train_loss)
    valid_losses.append(valid_loss)

```

```

Epoch: 1      Training Loss: 1.554870      Validation Loss: 1.195292
Validation loss decreased (inf --> 1.195292). Saving model ...
Epoch: 2      Training Loss: 1.061604      Validation Loss: 0.942772
Validation loss decreased (1.195292 --> 0.942772). Saving model ...
Epoch: 3      Training Loss: 0.855026      Validation Loss: 0.898186
Validation loss decreased (0.942772 --> 0.898186). Saving model ...
Epoch: 4      Training Loss: 0.721844      Validation Loss: 0.802594
Validation loss decreased (0.898186 --> 0.802594). Saving model ...
Epoch: 5      Training Loss: 0.624167      Validation Loss: 0.833805
Epoch: 6      Training Loss: 0.540125      Validation Loss: 0.847034
Epoch: 7      Training Loss: 0.477229      Validation Loss: 0.885064
Epoch: 8      Training Loss: 0.431222      Validation Loss: 0.974479
Epoch: 9      Training Loss: 0.396895      Validation Loss: 1.023753
Epoch: 10     Training Loss: 0.406966      Validation Loss: 1.029629

```

train_losses

```
[1.5548700734376908,
1.0616043160259723,
0.8550256190896034,
0.7218438929654658,
0.6241673525348306,
0.5401253384537995,
0.4772286603767425,
0.43122237718850376,
0.39689478159556163,
0.40696643604617566]
```

valid_losses

```
[1.195291963338852,
0.9427721498012542,
0.8981857757568359,
0.8025938410758973,
0.8338054661154747,
0.8470336732268333,
0.8850638046860695,
0.9744785497635603,
1.023752563506365,
1.0296290737092495]
```

```
# Plot the training and validation losses
import matplotlib.pyplot as plt
iepochs = len(valid_losses)
plt.figure(figsize=(10, 6))
plt.plot(range(1, iepochs + 1), train_losses, label='Training Loss')
plt.plot(range(1, iepochs + 1), valid_losses, label='Validation Loss')
```

```
plt.title('Training and Validation Losses')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```



✓ Load the Model with the Lowest Validation Loss

This is the model we will use for testing, which is the model we saved in the last step

```
model.load_state_dict(torch.load('model_trained_10ep_sgd.pt'))
```

```

<ipython-input-13-ef56f024e9ee>:1: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which
model.load_state_dict(torch.load('model_trained_10ep_sgd.pt'))
<All keys matched successfully>

```

✓ Test the Trained Network

Test your trained model on previously unseen data! Remember we have downloaded `train_data` and `test_data`. We will use `test_data` through `test_loader`.

A "good" result will be a CNN that gets around 70% (or more, try your best!) accuracy on these test images.

The following is working code, but you are encouraged to make your own adjustments and enhance the implementation.

```

# track test loss
test_loss = 0.0
class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))

model.eval()
# iterate over test data
for batch_idx, (data, target) in enumerate(test_loader):
    # move tensors to GPU if CUDA is available
    if train_on_gpu:
        data, target = data.cuda(), target.cuda()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the batch loss
    loss = criterion(output, target)
    # update test loss
    test_loss += loss.item()*data.size(0)
    # convert output probabilities to predicted class
    _, pred = torch.max(output, 1)
    # compare predictions to true label
    correct_tensor = pred.eq(target.data.view_as(pred))
    correct = np.squeeze(correct_tensor.numpy()) if not train_on_gpu else np.squeeze(correct_tensor.cpu().numpy())
    # calculate test accuracy for each object class
    for i in range(batch_size):
        label = target.data[i]
        class_correct[label] += correct[i].item()
        class_total[label] += 1

# average test loss
test_loss = test_loss/len(test_loader.dataset)
print('Test Loss: {:.6f}\n'.format(test_loss))

for i in range(10):
    if class_total[i] > 0:
        print('Test Accuracy of %5s: %2d%% (%2d/%2d)' % (
            classes[i], 100 * class_correct[i] / class_total[i],
            np.sum(class_correct[i]), np.sum(class_total[i])))
    else:
        print('Test Accuracy of %5s: N/A (no training examples)' % (classes[i]))

print('\nTest Accuracy (Overall): %2d%% (%2d/%2d)' % (
    100. * np.sum(class_correct) / np.sum(class_total),
    np.sum(class_correct), np.sum(class_total)))

```

```

Test Loss: 0.820112

```

```

Test Accuracy of airplane: 82% (828/1000)
Test Accuracy of automobile: 85% (852/1000)
Test Accuracy of bird: 70% (701/1000)
Test Accuracy of cat: 43% (437/1000)
Test Accuracy of deer: 71% (713/1000)
Test Accuracy of dog: 56% (568/1000)
Test Accuracy of frog: 79% (791/1000)
Test Accuracy of horse: 75% (753/1000)
Test Accuracy of ship: 78% (788/1000)
Test Accuracy of truck: 78% (780/1000)

```

```

Test Accuracy (Overall): 72% (7211/10000)

```


Visualize Sample Test Results

The following is working code, but you are encouraged to make your own adjustments and enhance the visualization.

```
# obtain one batch of test images
dataiter = iter(test_loader)
images, labels = next(dataiter)


# move model inputs to cuda, if GPU available
if train_on_gpu:
    images = images.cuda()

# get sample outputs
output = model(images)

# convert output probabilities to predicted class
_, preds_tensor = torch.max(output, 1)

# move the predictions to CPU if they're on the GPU and then convert to numpy
preds = np.squeeze(preds_tensor.cpu().numpy()) if not train_on_gpu else np.squeeze(preds_tensor.cpu().numpy())

# plot the images in the batch, along with predicted and true labels
fig = plt.figure(figsize=(25, 4))
for idx in np.arange(20):
    ax = fig.add_subplot(2, int(20/2), idx+1, xticks=[], yticks=[])
    imshow(images[idx].cpu()) # move the image to CPU for plotting if necessary
    ax.set_title("{} ({}).format(classes[preds[idx]], classes[labels[idx]]),
                color=("green" if preds[idx] == labels[idx].item() else "red"))
```



```
# obtain one batch of test images
dataiter = iter(test_loader)
images, labels = next(dataiter)
images.numpy()

# move model inputs to cuda, if GPU available
if train_on_gpu:
    images = images.cuda()

# get sample outputs
output = model(images)
# convert output probabilities to predicted class
_, preds_tensor = torch.max(output, 1)
preds = np.squeeze(preds_tensor.numpy()) if not train_on_gpu else np.squeeze(preds_tensor.cpu().numpy())

# plot the images in the batch, along with predicted and true labels
fig = plt.figure(figsize=(25, 4))
for idx in np.arange(20):
    ax = fig.add_subplot(2, int(20/2), idx+1, xticks=[], yticks=[])
    imshow(images[idx])
    ax.set_title("{} ({}).format(classes[preds[idx]], classes[labels[idx]]),
                color=("green" if preds[idx]==labels[idx].item() else "red"))
```

Start coding or [generate](#) with AI.

